

BDPA - Assignment 1

Hippolyte JACOMET - github.com/hippolytej/BDPA-HW

1 HADOOP SETUP

Here's my Hadoop setup, didn't change the defaults.

```
[cloudera@quickstart ~]$ hadoop version
Hadoop 2.6.0-cdh5.8.0
Subversion http://github.com/cloudera/hadoop -r
57e7b8556919574d517e874abfb7ebe31a366c2b
Compiled by jenkins on 2016-06-16T19:38Z
Compiled with protoc 2.5.0
From source with checksum 9e99ecd28376acfd5f78c325dd939fed
This command was run using /usr/lib/hadoop/hadoop-common-2.6.0-cdh5.8.0.jar
[cloudera@quickstart ~]$ hadoop checknative -a
17/02/16 11:55:57 INFO bzip2.Bzip2Factory: Successfully loaded &
initialized native-bzip2 library system-native
17/02/16 11:55:57 INFO zlib.ZlibFactory: Successfully loaded & initialized
native-zlib library
Native library checking:
hadoop: true /usr/lib/hadoop/lib/native/libhadoop.so.1.0.0
zlib: true /lib64/libz.so.1
snappy: true /usr/lib/hadoop/lib/native/libsnappy.so.1
lz4: true revision:10301
bzip2: true /lib64/libbz2.so.1
openssl: true /usr/lib64/libcrypto.so
```

As for my virtual machine:

- Memory 4096 Mo;
- 1 Cores 100% allocated (later added one for inverted index, was way too slow);
- 24 Mo of Graphics Memory (default was also too light).

2 IMPLEMENTATIONS

2.1 STOP WORDS

2.1.1 Main

Based on the wordcount main of the tutorial found at [wiki.apache.org](https://wiki.apache.org/hadoop/WordCount).

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = new Job(conf, "StopWords");
    job.setJarByClass(StopWords.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    job.setMapperClass(Map.class);
    // job.setCombinerClass(Combiner.class);
    job.setReducerClass(Reduce.class);
    job.setNumReduceTasks(10);
    // job.setNumReduceTasks(50);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.setInputPaths(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    // conf.set("mapreduce.map.output.compress", "true");
    // conf.set("mapreduce.map.output.compress.codec", "org.apache.ha-
doop.io.compress.SnappyCodec");

    job.waitForCompletion(true);
}

```

2.1.2 Mapper

Didn't change much the code used in the wordcount tutorial, mainly personalized the tokenizer to avoid inappropriate "words":

```

public static class Map extends Mapper<LongWritable, Text, Text,
IntWritable> {
    private final static IntWritable ONE = new IntWritable(1);
    private Text word = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString().toLowerCase();
        StringTokenizer tokenizer = new StringTokenizer(line, "
\t\n\r\f,.;?![]{}'\\"() &<>~_-1234567890#$%^%/\@\\\'`=+|"");
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            context.write(word, ONE);
        }
    }
}

```

2.1.3 Combiner

Had to create a custom combiner as I could not simply use the reducer class because I chose to write only the words and not the count in the reducer (cf. next section)

```

    public static class Combiner extends Reducer<Text,IntWritable,
Text,IntWritable> {
        @Override
        public void reduce(Text key, Iterable<IntWritable> values, Context
context) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }
}

```

2.1.4 Reducer

Simply added an if condition to the wordcount reducer, but also chose to not write the count which is useless: I changed the output to a NullWritable.

```

    public static class Reduce extends Reducer<Text, IntWritable, Text,
NullWritable> {
        @Override
        public void reduce(Text key, Iterable<IntWritable> values, Context
context)
            throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            if (sum > 4000){
                context.write(key, NullWritable.get());
            }
        }
    }
}

```

2.2 INVERTED INDEX

I detail here only the code of the inverted index with frequencies, which is more complete and features the counter. Code for the basic inverted index one can be found in my GitHub repository.

2.2.1 Main & counter

The main is quite similar to the one from stopwords apart from the counter which is saved in a file on hdfs.

```

static enum CustomCounters {UNIQUEWORDS}
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = new Job(conf, "InvertedIndex");
    job.setJarByClass(InvertedIndex.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    job.setMapperClass(Map.class);
    job.setCombinerClass(Combiner.class);
    job.setReducerClass(Reduce.class);
    job.setNumReduceTasks(10);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.setInputPaths(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    conf.set("mapreduce.map.output.compress", "true");
    conf.set("mapreduce.map.output.compress.codec", "org.apache.ha-
doop.io.compress.SnappyCodec");

    job.waitForCompletion(true);
    Counter counter = job.getCounters().findCounter(Cus-
tomCounters.UNIQUEWORDS);

    FileSystem hdfs = FileSystem.get(URI.create("count"), conf);
    Path file = new Path("counter.txt");
    if ( hdfs.exists( file )) { hdfs.delete( file, true ); }
    OutputStream os = hdfs.create(file);
    BufferedWriter br = new BufferedWriter( new OutputStreamWriter( os,
"UTF-8" ) );
    br.write("Unique words in a single file = " + counter.getValue());
    br.close();
    hdfs.close();
}

```

2.2.2 Mapper

Based on the same mapper as the one used for Stopwords, only changed the values output to Text, and wrote the file names extracted from the context variable.

```

    public static class Map extends Mapper<LongWritable, Text, Text, Text>
    {
        private Text word = new Text();

        @Override
        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {

            FileSplit split = (FileSplit) context.getInputSplit();
            String filename = split.getPath().getName().toString();

            String line = value.toString().toLowerCase();
            StringTokenizer tokenizer = new StringTokenizer(line, "
\t\n\r\f,.;?![]{}'\\" ( ) &<>~_ -12345677890#$%^%/\@\\\'`=+|"");
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, new Text(filename));
            }
        }
    }
}

```

2.2.3 Combiner

Had to create a custom combiner which slightly differs from the reducer. A hashmap is created as an intermediate index, saved as suggested in the form of doc1#count1,doc2#count2,... so as to later be easily processed by the reducers.

```

    public static class Combiner extends Reducer<Text,Text, Text,Text> {
        @Override
        public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
            HashMap<String, Integer> countDooku = new HashMap<String, Integer>();

            for (Text val : values) {
                if(countDooku.containsKey(val.toString())){
                    countDooku.put(val.toString(), countDooku.get(val.toString()) + 1);
                }else{
                    countDooku.put(val.toString(), 1);
                }
            }

            String filesFrequency = new String();
            for (String fileName: countDooku.keySet()){
                String freq = countDooku.get(fileName).toString();
                filesFrequency = filesFrequency + fileName + "#" + freq +
", ";
            }

            filesFrequency = filesFrequency.substring(0, filesFrequency.length()-1);

            Text index = new Text();
            index.set(filesFrequency.toString());
            context.write(key, index);
        }
    }
}

```

2.2.4 Reducer

I chose to rule out the stop words only in the reducer as it is where the number of comparison with stopwords.csv would be minimal as the keys have been concatenated. Hence, the reducer starts with importing the .csv file, and checking if the key is a stop word. I used code for both local and hdfs files to be able to debug in standalone mode.

In the reduce method, a hashmap is used in the same fashion as the combiner, except the input is slightly different and must be split along “,” and “#”.

Finally, the counter is incremented if a key has only one document name in its output value, that is, simply if no comma can be found in it.

```

    public static class Reduce extends Reducer<Text, Text, Text, Text> {
        String stopwords = new String();
        public void setup(Context context) throws IOException,
InterruptedException {
            // Check if key is a stop word -----
            // Test with local file for standalone mode
            // File file = new File("stopwords.csv");
            // Scanner sw = new Scanner(file);
            // With DHFS file
            Path pt=new Path("stopwords.csv");
            FileSystem fs = FileSystem.get(new Configuration());
            Scanner sw=new Scanner(fs.open(pt));

            while (sw.hasNext()){
                stopwords = stopwords + " " + sw.next().toString();
            }
            sw.close();
        }
        @Override
        public void reduce(Text key, Iterable<Text> values, Context
context)
            throws IOException, InterruptedException {
            // Create index and count frequencies -----
            if (!stopwords.contains(key.toString())){
                HashMap<String, Integer> countDooku = new HashMap<String,
Integer>();
                for (Text val : values) {
                    for (String token: val.toString().split(",")) {
                        String[] couple = token.split("#");
                        if(countDooku.containsKey(couple[0])){
                            countDooku.put(couple[0],
countDooku.get(couple[0]) + Integer.parseInt(couple[1]));
                        }
                        else{
                            countDooku.put(couple[0],
Integer.parseInt(couple[1]));
                        }
                    }
                }
                String finalFilesFrequency = new String();
                for (String fileName: countDooku.keySet()){
                    String freq = countDooku.get(fileName).toString();
                    finalFilesFrequency = finalFilesFrequency + fileName +
"#" + freq + ", ";
                }
                finalFilesFrequency = finalFilesFrequency.substring(0,
finalFilesFrequency.length()-2);
                Text finalIndex = new Text();
                finalIndex.set(finalFilesFrequency);
                context.write(key, finalIndex);
            }
            // Here comes the counter...
            if (!finalFilesFrequency.contains(",")){
                context.getCounter(CustomCounters.UNIQUEWORDS).increment(1);
            }
        }
    }
}

```

3 DATA SET

I made the simplification of excluding all numbers and not count them as words, even though there may be dates, and more generally many special characters that are less controversial.

4 TEST SCENARIOS

The procedure for running tests was the following: update code on eclipse accordingly to test (adjust the number of reducers...), export as jar, and run with Hadoop.

Results were concatenated and saved on hdfs using:

```
hadoop fs -cat output/part\* | hadoop fs -put - name.text
```

And then locally downloaded with:

```
Hadoop -get name.text
```

4.1 STOPWORDS WITH 10 REDUCERS

Execution time : 2'19''

Job Overview			
Job Name: StopWords			
User Name: cloudera			
Queue: root.cloudera			
State: SUCCEEDED			
Uberized: false			
Submitted: Wed Feb 15 06:31:29 PST 2017			
Started: Wed Feb 15 06:31:38 PST 2017			
Finished: Wed Feb 15 06:33:58 PST 2017			
Elapsed: 2mins, 19sec			
Diagnostics:			
Average Map Time 35sec			
Average Shuffle Time 36sec			
Average Merge Time 5sec			
Average Reduce Time 9sec			

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 06:31:32 PST 2017	quickstart.cloudera:8042	logs

Task Type	Total		Complete	
Map	3		3	
Reduce	10		10	
Attempt Type	Failed	Killed	Successful	
Maps	<u>0</u>	<u>0</u>	<u>3</u>	
Reduces	<u>0</u>	<u>3</u>	<u>10</u>	

17/02/15 06:34:00 INFO mapreduce.Job: Job job_1485279029005_0017 completed successfully

17/02/15 06:34:00 INFO mapreduce.Job: Counters: 50

File System Counters

FILE: Number of bytes read=51484994
FILE: Number of bytes written=104487032
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=26058246
HDFS: Number of bytes written=611
HDFS: Number of read operations=39
HDFS: Number of large read operations=0
HDFS: Number of write operations=20

Job Counters

Killed reduce tasks=3
Launched map tasks=3
Launched reduce tasks=13
Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)=105366
Total time spent by all reduces in occupied slots

(ms)=551481

Total time spent by all map tasks (ms)=105366
Total time spent by all reduce tasks (ms)=551481
Total vcore-seconds taken by all map tasks=105366
Total vcore-seconds taken by all reduce tasks=551481
Total megabyte-seconds taken by all map tasks=107894784
Total megabyte-seconds taken by all reduce tasks=564716544

Map-Reduce Framework

Map input records=507535
Map output records=4603237
Map output bytes=42278460
Map output materialized bytes=51485114
Input split bytes=381
Combine input records=0
Combine output records=0
Reduce input groups=56726
Reduce shuffle bytes=51485114
Reduce input records=4603237
Reduce output records=139
Spilled Records=9206474
Shuffled Maps =30
Failed Shuffles=0
Merged Map outputs=30
GC time elapsed (ms)=3699
CPU time spent (ms)=41830
Physical memory (bytes) snapshot=1748418560
Virtual memory (bytes) snapshot=19566936064
Total committed heap usage (bytes)=1104359424

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=26057865

File Output Format Counters

Bytes Written=611

4.2 STOPWORDS WITH 10 REDUCERS AND COMBINER

Execution time : 1'27'' A significant improvement is observed here, which is what we expect of the combiner : lighter files are sent to the reducer.

Job Overview			
Job Name:	StopWords		
User Name:	cloudera		
Queue:	root.cloudera		
State:	SUCCEEDED		
Uberized:	false		
Submitted:	Thu Feb 16 14:47:15 PST 2017		
Started:	Thu Feb 16 14:47:23 PST 2017		
Finished:	Thu Feb 16 14:48:42 PST 2017		
Elapsed:	1mins, 19sec		
Diagnostics:			
Average Map Time	29sec		
Average Shuffle Time	19sec		
Average Merge Time	0sec		
Average Reduce Time	2sec		

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Thu Feb 16 14:47:18 PST 2017	quickstart.cloudera:8042	logs

Task Type	Total		Complete	
Map	3		3	
Reduce	10		10	
Attempt Type	Failed	Killed	Successful	
Maps	0	0	3	
Reduces	0	0	10	

17/02/16 14:48:44 INFO mapreduce.Job: Job job_1487258453206_0005 completed successfully

17/02/16 14:48:44 INFO mapreduce.Job: Counters: 51

File System Counters

FILE: Number of bytes read=1240085
FILE: Number of bytes written=3999308
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=26058246
HDFS: Number of bytes written=611
HDFS: Number of read operations=39
HDFS: Number of large read operations=0
HDFS: Number of write operations=20

Job Counters

Killed map tasks=1
Killed reduce tasks=2
Launched map tasks=3
Launched reduce tasks=10
Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)=89028
Total time spent by all reduces in occupied slots (ms)=228144
Total time spent by all map tasks (ms)=89028
Total time spent by all reduce tasks (ms)=228144
Total vcore-seconds taken by all map tasks=89028
Total vcore-seconds taken by all reduce tasks=228144
Total megabyte-seconds taken by all map tasks=91164672
Total megabyte-seconds taken by all reduce tasks=233619456

Map-Reduce Framework

Map input records=507535
Map output records=4603237
Map output bytes=42278460
Map output materialized bytes=1240205
Input split bytes=381
Combine input records=4603237
Combine output records=85622
Reduce input groups=56726
Reduce shuffle bytes=1240205
Reduce input records=85622
Reduce output records=139
Spilled Records=171244
Shuffled Maps =30
Failed Shuffles=0
Merged Map outputs=30
GC time elapsed (ms)=2149
CPU time spent (ms)=38070
Physical memory (bytes) snapshot=2589450240
Virtual memory (bytes) snapshot=20322811904
Total committed heap usage (bytes)=1916796928

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=26057865

File Output Format Counters

Bytes Written=611

4.3 STOPWORDS WITH 10 REDUCERS, COMBINER, AND COMPRESSION

Execution time 1'11''. A slight improvement is observed but nothing major, maybe because the files are not big enough or the fact that it is only a pseudo distributed version : little to no time is gain in transfer, but time is lost in compressing and decompressing.

Job Overview			
Job Name:	StopWords		
User Name:	cloudera		
Queue:	root.cloudera		
State:	SUCCEEDED		
Uberized:	false		
Submitted:	Thu Feb 16 14:56:00 PST 2017		
Started:	Thu Feb 16 14:56:06 PST 2017		
Finished:	Thu Feb 16 14:57:17 PST 2017		
Elapsed:	1mins, 11sec		
Diagnostics:			
Average Map Time	23sec		
Average Shuffle Time	18sec		
Average Merge Time	0sec		
Average Reduce Time	3sec		

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Thu Feb 16 14:56:01 PST 2017	quickstart.cloudera:8042	logs

Task Type	Total		Complete	
Map	3		3	
Reduce	10		10	
Attempt Type	Failed	Killed	Successful	
Maps	0	0	3	
Reduces	0	1	10	

17/02/15 07:10:29 INFO mapreduce.Job: Job job_1485279029005_0019 completed successfully

17/02/15 07:10:29 INFO mapreduce.Job: Counters: 50

File System Counters

FILE: Number of bytes read=1240085
FILE: Number of bytes written=3999373
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=26058246
HDFS: Number of bytes written=611
HDFS: Number of read operations=39
HDFS: Number of large read operations=0
HDFS: Number of write operations=20

Job Counters

Killed reduce tasks=2
Launched map tasks=3
Launched reduce tasks=12
Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)=101483
Total time spent by all reduces in occupied slots (ms)=419584
Total time spent by all map tasks (ms)=101483
Total time spent by all reduce tasks (ms)=419584
Total vcore-seconds taken by all map tasks=101483
Total vcore-seconds taken by all reduce tasks=419584
Total megabyte-seconds taken by all map tasks=103918592
Total megabyte-seconds taken by all reduce tasks=429654016

Map-Reduce Framework

Map input records=507535
Map output records=4603237
Map output bytes=42278460
Map output materialized bytes=1240205
Input split bytes=381
Combine input records=4603237
Combine output records=85622
Reduce input groups=56726
Reduce shuffle bytes=1240205
Reduce input records=85622
Reduce output records=139
Spilled Records=171244
Shuffled Maps =30
Failed Shuffles=0
Merged Map outputs=30
GC time elapsed (ms)=2974
CPU time spent (ms)=31130
Physical memory (bytes) snapshot=1774018560
Virtual memory (bytes) snapshot=19566510080
Total committed heap usage (bytes)=1104359424

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=26057865

File Output Format Counters

Bytes Written=611

17/02/15 07:22:06 INFO mapreduce.Job: Job job_1485279029005_0020 completed successfully

17/02/15 07:22:06 INFO mapreduce.Job: Counters: 50

File System Counters

FILE: Number of bytes read=1240325
FILE: Number of bytes written=8674853
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=26058246
HDFS: Number of bytes written=611
HDFS: Number of read operations=159
HDFS: Number of large read operations=0
HDFS: Number of write operations=100

Job Counters

Killed reduce tasks=2
Launched map tasks=3
Launched reduce tasks=50
Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)=102506
Total time spent by all reduces in occupied slots (ms)=1998484
Total time spent by all map tasks (ms)=102506
Total time spent by all reduce tasks (ms)=1998484
Total vcore-seconds taken by all map tasks=102506
Total vcore-seconds taken by all reduce tasks=1998484
Total megabyte-seconds taken by all map tasks=104966144
Total megabyte-seconds taken by all reduce tasks=2046447616

Map-Reduce Framework

Map input records=507535
Map output records=4603237
Map output bytes=42278460
Map output materialized bytes=1240925
Input split bytes=381
Combine input records=4603237
Combine output records=85622
Reduce input groups=56726
Reduce shuffle bytes=1240925
Reduce input records=85622
Reduce output records=139
Spilled Records=171244
Shuffled Maps =150
Failed Shuffles=0
Merged Map outputs=150
GC time elapsed (ms)=11562
CPU time spent (ms)=71060
Physical memory (bytes) snapshot=6639992832
Virtual memory (bytes) snapshot=79839502336
Total committed heap usage (bytes)=3534434304

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=26057865

File Output Format Counters

Bytes Written=611

4.5 INVERTED INDEX

Here are the performances of my inverted index. The counter `Reduce output records=55643` shows the number of unique words there are, excluding stopwords.

Job Overview	
Job Name:	InvertedIndex
User Name:	cloudera
Queue:	root.cloudera
State:	SUCCEEDED
Uberized:	false
Submitted:	Thu Feb 16 07:56:52 PST 2017
Started:	Thu Feb 16 07:57:06 PST 2017
Finished:	Thu Feb 16 08:09:56 PST 2017
Elapsed:	12mins, 49sec
Diagnostics:	
Average Map Time	31sec
Average Shuffle Time	18sec
Average Merge Time	0sec
Average Reduce Time	5mins, 56sec

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Thu Feb 16 07:57:00 PST 2017	quickstart.cloudera:8042	logs

Task Type	Total		Complete	
Map	3		3	
Reduce	10		10	
Attempt Type	Failed		Killed	Successful
Maps	0	0		3
Reduces	0	1		10

17/02/16 08:09:58 INFO mapreduce.Job: Job job_1487258453206_0001 completed successfully

17/02/16 08:09:59 INFO mapreduce.Job: Counters: 51

File System Counters

FILE: Number of bytes read=3761127
FILE: Number of bytes written=7555863
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=60141048
HDFS: Number of bytes written=1622956
HDFS: Number of read operations=55821
HDFS: Number of large read operations=0
HDFS: Number of write operations=20

Job Counters

Killed reduce tasks=1
Launched map tasks=3
Launched reduce tasks=11
Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)=93941
Total time spent by all reduces in occupied slots (ms)=4064026
Total time spent by all map tasks (ms)=93941
Total time spent by all reduce tasks (ms)=4064026
Total vcore-seconds taken by all map tasks=93941
Total vcore-seconds taken by all reduce tasks=4064026
Total megabyte-seconds taken by all map tasks=96195584
Total megabyte-seconds taken by all reduce tasks=4161562624

Map-Reduce Framework

Map input records=507535
Map output records=4591847
Map output bytes=74198099
Map output materialized bytes=2313290
Input split bytes=381
Combine input records=4591847
Combine output records=97518
Reduce input groups=55782
Reduce shuffle bytes=2313290
Reduce input records=97518
Reduce output records=55643
Spilled Records=253971
Shuffled Maps =30
Failed Shuffles=0
Merged Map outputs=30
GC time elapsed (ms)=36668
CPU time spent (ms)=1184860
Physical memory (bytes) snapshot=3341504512
Virtual memory (bytes) snapshot=20355305472
Total committed heap usage (bytes)=2181562368

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

invertedindex.InvertedIndex\$CustomCounters

UNIQUEWORDS=35726

File Input Format Counters

Bytes Read=26057865

File Output Format Counters

Bytes Written=1622956

5 CONCLUSION

This homework shows the influence of critical parameters of Hadoop implementation: number of reducers, combiners and compression, loading and saving files on hdfs, as well as the interesting features of counters.

SOURCES:

wiki.apache.org

Mining of Massive Datasets, Jure Leskovec, Anand Rajaraman, Jeff Ullman

Countless hours of web searches